

Z.ai Engineering Handbook

This document describes the engineering practices, deployment workflow, and on-call procedures used by the Z.ai platform team. It is intended for all engineers joining the team.

1. Team Structure

The platform team consists of 12 engineers split across three squads: Infrastructure (4 people), Application Platform (5 people), and Developer Experience (3 people). Each squad has one tech lead who reports to the VP of Engineering, Priya Sharma.

2. Deployment Workflow

All production deployments use the following workflow: pull request → automated CI tests → code review by two engineers → merge to main → auto-deploy to staging → smoke tests → manual promotion to production. The average deployment takes 18 minutes from merge to production.

Production deploys happen on Tuesdays and Thursdays between 10 AM and 2 PM Pacific Time. Hotfixes can be deployed at any time but require approval from the on-call engineer and the tech lead of the affected squad.

3. On-Call Rotation

The on-call rotation is weekly, running Monday to Monday. Each engineer is on-call approximately once every 6 weeks. The on-call engineer receives a \$400 stipend per week of on-call duty. Critical alerts are routed via PagerDuty and must be acknowledged within 5 minutes.

4. Infrastructure

Production infrastructure runs on AWS in the us-west-2 region. The stack includes: EKS for Kubernetes, RDS Postgres for relational data, ElastiCache Redis for caching, S3 for object storage, and CloudFront for CDN. Monthly infrastructure cost averages \$42,000.

5. SLA and Uptime

The platform SLA is 99.95% monthly uptime for the API tier. In Q3 2025, actual uptime was 99.97%. The largest source of downtime was a 23-minute outage on August 14 caused by a misconfigured load balancer.